

# major-final-5

August 17, 2025

```
[1]: import os
import numpy as np
import cv2
from sklearn.model_selection import train_test_split
from tensorflow.keras.preprocessing.image import array_to_img

# Classes
classes = ["Healthy", "Early_blight", "Late_blight", "Yellow_Leaf_Curl_Virus"]

# Base path
base_dir = "/content/tomato_dataset"
os.makedirs(base_dir+"/train", exist_ok=True)
os.makedirs(base_dir+"/val", exist_ok=True)

# Function to generate synthetic images
def generate_images(class_name, n=250):
    os.makedirs(f"{base_dir}/train/{class_name}", exist_ok=True)
    os.makedirs(f"{base_dir}/val/{class_name}", exist_ok=True)
    for i in range(n):
        # create synthetic image (random noise with slight patterns)
        img = np.zeros((224,224,3), dtype=np.uint8)
        cv2.putText(img, class_name[:2], (50,100),
                    cv2.FONT_HERSHEY_SIMPLEX, 2,
                    (np.random.randint(0,255),
                     np.random.randint(0,255),
                     np.random.randint(0,255)), 3)
        # split into train/val (80/20)
        if i < int(n*0.8):
            cv2.imwrite(f"{base_dir}/train/{class_name}/{class_name}_{i}.jpg", img)
        else:
            cv2.imwrite(f"{base_dir}/val/{class_name}/{class_name}_{i}.jpg", img)

# Generate 4 classes
for cls in classes:
    generate_images(cls, n=250)
```

```
print(" Synthetic dataset created with 1000 images (4 classes, 80/20 split)")
```

Synthetic dataset created with 1000 images (4 classes, 80/20 split)

```
[2]: from tensorflow.keras.preprocessing.image import ImageDataGenerator

train_dir = "/content/tomato_dataset/train"
val_dir = "/content/tomato_dataset/val"

datagen = ImageDataGenerator(rescale=1./255, rotation_range=20,
                             zoom_range=0.2, horizontal_flip=True)

train_data = datagen.flow_from_directory(train_dir, target_size=(224,224),
                                         batch_size=32, class_mode='categorical')

val_data = datagen.flow_from_directory(val_dir, target_size=(224,224),
                                         batch_size=32, class_mode='categorical')
```

Found 800 images belonging to 4 classes.

Found 200 images belonging to 4 classes.

```
[3]: import matplotlib.pyplot as plt
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout, GlobalAveragePooling2D
from tensorflow.keras.optimizers import Adam

def train_model(base_model, model_name, epochs=10, lr=1e-4):
    base_model.trainable = False # Freeze pretrained layers

    model = Sequential([
        base_model,
        GlobalAveragePooling2D(),
        Dense(128, activation='relu'),
        Dropout(0.3),
        Dense(train_data.num_classes, activation='softmax')
    ])

    model.compile(optimizer=Adam(learning_rate=lr),
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])

    history = model.fit(train_data, validation_data=val_data, epochs=epochs)

    # Accuracy plot
    plt.plot(history.history['accuracy'], label='Train Acc')
```

```

plt.plot(history.history['val_accuracy'], label='Val Acc')
plt.title(model_name + " Accuracy")
plt.legend()
plt.show()

val_loss, val_acc = model.evaluate(val_data)
print(f"{model_name} Final Accuracy: {val_acc*100:.2f}%")
return val_acc*100

```

[1]: from tensorflow.keras.applications import VGG19, MobileNetV2, ResNet50

```

vgg_acc = train_model(VGG19(weights='imagenet', include_top=False, □
    ↪input_shape=(224,224,3)), "VGG19")
mob_acc = train_model(MobileNetV2(weights='imagenet', include_top=False, □
    ↪input_shape=(224,224,3)), "MobileNetV2")
res_acc = train_model(ResNet50(weights='imagenet', include_top=False, □
    ↪input_shape=(224,224,3)), "ResNet50")

print("\n Model Comparison Results:")
print(f"VGG19 Accuracy : {vgg_acc:.2f}%")
print(f"MobileNetV2 Accuracy: {mob_acc:.2f}%")
print(f"ResNet50 Accuracy : {res_acc:.2f}%")

```

```

-----
NameError                                 Traceback (most recent call last)
/tmp/ipython-input-445742229.py in <cell line: 0>()
      1 from tensorflow.keras.applications import VGG19, MobileNetV2, ResNet50
      2
----> 3 vgg_acc = train_model(VGG19(weights='imagenet', include_top=False, □
    ↪input_shape=(224,224,3)), "VGG19")
      4 mob_acc = train_model(MobileNetV2(weights='imagenet', include_top=False, □
    ↪input_shape=(224,224,3)), "MobileNetV2")
      5 res_acc = train_model(ResNet50(weights='imagenet', include_top=False, □
    ↪input_shape=(224,224,3)), "ResNet50")

NameError: name 'train_model' is not defined

```

[3]: import matplotlib.pyplot as plt  
from tensorflow.keras.models import Sequential  
from tensorflow.keras.layers import Dense, Dropout, GlobalAveragePooling2D  
from tensorflow.keras.optimizers import Adam  
from tensorflow.keras.applications import VGG19, MobileNetV2, ResNet50  
from tensorflow.keras.preprocessing.image import ImageDataGenerator

```

train_dir = "/content/tomato_dataset/train"
val_dir = "/content/tomato_dataset/val"

```

```

datagen = ImageDataGenerator(rescale=1./255, rotation_range=20,
                             zoom_range=0.2, horizontal_flip=True)

train_data = datagen.flow_from_directory(train_dir, target_size=(224,224),
                                         batch_size=32, class_mode='categorical')

val_data = datagen.flow_from_directory(val_dir, target_size=(224,224),
                                         batch_size=32, class_mode='categorical')

def train_model(base_model, model_name, epochs=10, lr=1e-4):
    base_model.trainable = False # Freeze pretrained layers

    model = Sequential([
        base_model,
        GlobalAveragePooling2D(),
        Dense(128, activation='relu'),
        Dropout(0.3),
        Dense(train_data.num_classes, activation='softmax')
    ])

    model.compile(optimizer=Adam(learning_rate=lr),
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])

    history = model.fit(train_data, validation_data=val_data, epochs=epochs)

    # Accuracy plot
    plt.plot(history.history['accuracy'], label='Train Acc')
    plt.plot(history.history['val_accuracy'], label='Val Acc')
    plt.title(model_name + " Accuracy")
    plt.legend()
    plt.show()

    val_loss, val_acc = model.evaluate(val_data)
    print(f"{model_name} Final Accuracy: {val_acc*100:.2f}%")
    return val_acc*100

vgg_acc = train_model(VGG19(weights='imagenet', include_top=False, input_shape=(224,224,3)), "VGG19")
mob_acc = train_model(MobileNetV2(weights='imagenet', include_top=False, input_shape=(224,224,3)), "MobileNetV2")
res_acc = train_model(ResNet50(weights='imagenet', include_top=False, input_shape=(224,224,3)), "ResNet50")

print("\n Model Comparison Results:")

```

```

print(f"VGG19 Accuracy      : {vgg_acc:.2f}%")
print(f"MobileNetV2 Accuracy: {mob_acc:.2f}%")
print(f"ResNet50 Accuracy   : {res_acc:.2f}%")

```

```

-----
FileNotFoundError                         Traceback (most recent call last)
/tmp/ipython-input-561608046.py in <cell line: 0>()
      12                                         zoom_range=0.2, horizontal_flip=True)
      13
--> 14 train_data = datagen.flow_from_directory(train_dir,
      ↪target_size=(224,224),
      15                                         batch_size=32, ↪
      ↪class_mode='categorical')
      16

/usr/local/lib/python3.11/dist-packages/keras/src/legacy/preprocessing/image.py
  ↪in flow_from_directory(self, directory, target_size, color_mode, classes, ↪
  ↪class_mode, batch_size, shuffle, seed, save_to_dir, save_prefix, save_format, ↪
  ↪follow_links, subset, interpolation, keep_aspect_ratio)
  1136         keep_aspect_ratio=False,
  1137     ):
-> 1138         return DirectoryIterator(
  1139             directory,
  1140             self,

/usr/local/lib/python3.11/dist-packages/keras/src/legacy/preprocessing/image.py
  ↪in __init__(self, directory, image_data_generator, target_size, color_mode, ↪
  ↪classes, class_mode, batch_size, shuffle, seed, data_format, save_to_dir, ↪
  ↪save_prefix, save_format, follow_links, subset, interpolation, ↪
  ↪keep_aspect_ratio, dtype)
  451         if not classes:
  452             classes = []
--> 453         for subdir in sorted(os.listdir(directory)):
  454             if os.path.isdir(os.path.join(directory, subdir)):
  455                 classes.append(subdir)

FileNotFoundError: [Errno 2] No such file or directory: '/content/tomato_dataset' / ↪
      ↪train'

```

```

[4]: import os
import numpy as np
import cv2
from sklearn.model_selection import train_test_split
from tensorflow.keras.preprocessing.image import array_to_img

# Classes
classes = ["Healthy", "Early_blight", "Late_blight", "Yellow_Leaf_Curl_Virus"]

```

```

# Base path
base_dir = "/content/tomato_dataset"
os.makedirs(base_dir+"/train", exist_ok=True)
os.makedirs(base_dir+"/val", exist_ok=True)

# Function to generate synthetic images
def generate_images(class_name, n=250):
    os.makedirs(f"{base_dir}/train/{class_name}", exist_ok=True)
    os.makedirs(f"{base_dir}/val/{class_name}", exist_ok=True)
    for i in range(n):
        # create synthetic image (random noise with slight patterns)
        img = np.zeros((224,224,3), dtype=np.uint8)
        cv2.putText(img, class_name[:2], (50,100),
                    cv2.FONT_HERSHEY_SIMPLEX, 2,
                    (np.random.randint(0,255),
                     np.random.randint(0,255),
                     np.random.randint(0,255)), 3)
        # split into train/val (80/20)
        if i < int(n*0.8):
            cv2.imwrite(f"{base_dir}/train/{class_name}/{class_name}_{i}.jpg", img)
        else:
            cv2.imwrite(f"{base_dir}/val/{class_name}/{class_name}_{i}.jpg", img)

# Generate 4 classes
for cls in classes:
    generate_images(cls, n=250)

print(" Synthetic dataset created with 1000 images (4 classes, 80/20 split)")

import matplotlib.pyplot as plt
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout, GlobalAveragePooling2D
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.applications import VGG19, MobileNetV2, ResNet50
from tensorflow.keras.preprocessing.image import ImageDataGenerator

train_dir = "/content/tomato_dataset/train"
val_dir = "/content/tomato_dataset/val"

datagen = ImageDataGenerator(rescale=1./255, rotation_range=20,
                            zoom_range=0.2, horizontal_flip=True)

train_data = datagen.flow_from_directory(train_dir, target_size=(224,224),
                                         batch_size=32, shuffle=True)

```

```

        batch_size=32, ↵
↳ class_mode='categorical')

val_data = datagen.flow_from_directory(val_dir, target_size=(224,224),
                                         batch_size=32, class_mode='categorical')

def train_model(base_model, model_name, epochs=10, lr=1e-4):
    base_model.trainable = False # Freeze pretrained layers

    model = Sequential([
        base_model,
        GlobalAveragePooling2D(),
        Dense(128, activation='relu'),
        Dropout(0.3),
        Dense(train_data.num_classes, activation='softmax')
    ])

    model.compile(optimizer=Adam(learning_rate=lr),
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])

    history = model.fit(train_data, validation_data=val_data, epochs=epochs)

    # Accuracy plot
    plt.plot(history.history['accuracy'], label='Train Acc')
    plt.plot(history.history['val_accuracy'], label='Val Acc')
    plt.title(model_name + " Accuracy")
    plt.legend()
    plt.show()

    val_loss, val_acc = model.evaluate(val_data)
    print(f"[{model_name}] Final Accuracy: {val_acc*100:.2f}%")
    return val_acc*100


vgg_acc = train_model(VGG19(weights='imagenet', include_top=False, ↵
↳ input_shape=(224,224,3)), "VGG19")
mob_acc = train_model(MobileNetV2(weights='imagenet', include_top=False, ↵
↳ input_shape=(224,224,3)), "MobileNetV2")
res_acc = train_model(ResNet50(weights='imagenet', include_top=False, ↵
↳ input_shape=(224,224,3)), "ResNet50")

print("\n Model Comparison Results:")
print(f"VGG19 Accuracy : {vgg_acc:.2f}%")
print(f"MobileNetV2 Accuracy: {mob_acc:.2f}%")
print(f"ResNet50 Accuracy : {res_acc:.2f}%")

```

Synthetic dataset created with 1000 images (4 classes, 80/20 split)

```
Found 800 images belonging to 4 classes.  
Found 200 images belonging to 4 classes.  
  
/usr/local/lib/python3.11/dist-  
packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:121:  
UserWarning: Your `PyDataset` class should call `super().__init__(**kwargs)` in  
its constructor. `**kwargs` can include `workers`, `use_multiprocessing`,  
`max_queue_size`. Do not pass these arguments to `fit()`, as they will be  
ignored.  
    self._warn_if_super_not_called()  
  
Epoch 1/10  
25/25          36s 807ms/step -  
accuracy: 0.2268 - loss: 1.5303 - val_accuracy: 0.2500 - val_loss: 1.3855  
Epoch 2/10  
25/25          13s 511ms/step -  
accuracy: 0.2528 - loss: 1.4417 - val_accuracy: 0.4400 - val_loss: 1.3745  
Epoch 3/10  
25/25          13s 508ms/step -  
accuracy: 0.2510 - loss: 1.4273 - val_accuracy: 0.5500 - val_loss: 1.3614  
Epoch 4/10  
25/25          13s 503ms/step -  
accuracy: 0.2802 - loss: 1.4046 - val_accuracy: 0.6000 - val_loss: 1.3475  
Epoch 5/10  
25/25          13s 503ms/step -  
accuracy: 0.2804 - loss: 1.4144 - val_accuracy: 0.5950 - val_loss: 1.3354  
Epoch 6/10  
25/25          13s 505ms/step -  
accuracy: 0.3329 - loss: 1.3835 - val_accuracy: 0.6450 - val_loss: 1.3252  
Epoch 7/10  
25/25          13s 502ms/step -  
accuracy: 0.3720 - loss: 1.3502 - val_accuracy: 0.7400 - val_loss: 1.3085  
Epoch 8/10  
25/25          13s 506ms/step -  
accuracy: 0.3623 - loss: 1.3448 - val_accuracy: 0.7350 - val_loss: 1.3021  
Epoch 9/10  
25/25          13s 507ms/step -  
accuracy: 0.3547 - loss: 1.3364 - val_accuracy: 0.6850 - val_loss: 1.2879  
Epoch 10/10  
25/25          13s 507ms/step -  
accuracy: 0.3728 - loss: 1.3182 - val_accuracy: 0.8300 - val_loss: 1.2752
```



```


7/7          3s 331ms/step -
accuracy: 0.7906 - loss: 1.2804
VGG19 Final Accuracy: 81.00%
Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/mobilenet_v2/mobilenet_v2_weights_tf_dim_ordering_tf_kernels_1.0_224_no_top.h5
9406464/9406464           0s
0us/step
Epoch 1/10
25/25          27s 719ms/step -
accuracy: 0.2405 - loss: 1.6924 - val_accuracy: 0.8250 - val_loss: 1.0042
Epoch 2/10
25/25          11s 428ms/step -
accuracy: 0.5940 - loss: 0.9972 - val_accuracy: 0.9250 - val_loss: 0.6366
Epoch 3/10
25/25          13s 516ms/step -
accuracy: 0.8104 - loss: 0.6827 - val_accuracy: 0.9800 - val_loss: 0.4556
Epoch 4/10
25/25          11s 425ms/step -
accuracy: 0.8922 - loss: 0.4862 - val_accuracy: 0.9850 - val_loss: 0.3684
Epoch 5/10

```

```

25/25      10s 415ms/step -
accuracy: 0.9507 - loss: 0.3599 - val_accuracy: 0.9800 - val_loss: 0.2477
Epoch 6/10
25/25      21s 421ms/step -
accuracy: 0.9522 - loss: 0.3143 - val_accuracy: 0.9900 - val_loss: 0.1822
Epoch 7/10
25/25      11s 424ms/step -
accuracy: 0.9708 - loss: 0.2253 - val_accuracy: 1.0000 - val_loss: 0.1634
Epoch 8/10
25/25      11s 427ms/step -
accuracy: 0.9687 - loss: 0.2367 - val_accuracy: 1.0000 - val_loss: 0.1471
Epoch 9/10
25/25      11s 424ms/step -
accuracy: 0.9877 - loss: 0.1723 - val_accuracy: 0.9950 - val_loss: 0.1263
Epoch 10/10
25/25      11s 425ms/step -
accuracy: 0.9923 - loss: 0.1316 - val_accuracy: 0.9950 - val_loss: 0.1149

```




```

7/7      2s 268ms/step -
accuracy: 0.9905 - loss: 0.1064
MobileNetV2 Final Accuracy: 98.50%
Downloading data from https://storage.googleapis.com/tensorflow/keras-

```

```
applications/resnet/resnet50_weights_tf_dim_ordering_tf_kernels_notop.h5
94765736/94765736          0s
0us/step
Epoch 1/10
25/25           32s 770ms/step -
accuracy: 0.1681 - loss: 1.5720 - val_accuracy: 0.2600 - val_loss: 1.3865
Epoch 2/10
25/25           11s 446ms/step -
accuracy: 0.2938 - loss: 1.4381 - val_accuracy: 0.2700 - val_loss: 1.3823
Epoch 3/10
25/25           11s 447ms/step -
accuracy: 0.2110 - loss: 1.5099 - val_accuracy: 0.2900 - val_loss: 1.3782
Epoch 4/10
25/25           13s 546ms/step -
accuracy: 0.2877 - loss: 1.4050 - val_accuracy: 0.2750 - val_loss: 1.3750
Epoch 5/10
25/25           11s 450ms/step -
accuracy: 0.2707 - loss: 1.4282 - val_accuracy: 0.2650 - val_loss: 1.3753
Epoch 6/10
25/25           13s 543ms/step -
accuracy: 0.2693 - loss: 1.3858 - val_accuracy: 0.3750 - val_loss: 1.3686
Epoch 7/10
25/25           11s 450ms/step -
accuracy: 0.2541 - loss: 1.4052 - val_accuracy: 0.4500 - val_loss: 1.3665
Epoch 8/10
25/25           11s 449ms/step -
accuracy: 0.2974 - loss: 1.3839 - val_accuracy: 0.3350 - val_loss: 1.3662
Epoch 9/10
25/25           11s 444ms/step -
accuracy: 0.2741 - loss: 1.3909 - val_accuracy: 0.3950 - val_loss: 1.3628
Epoch 10/10
25/25           11s 427ms/step -
accuracy: 0.3112 - loss: 1.3902 - val_accuracy: 0.3800 - val_loss: 1.3622
```



7/7                    2s 281ms/step -  
accuracy: 0.3883 - loss: 1.3577  
ResNet50 Final Accuracy: 40.00%

Model Comparison Results:  
VGG19 Accuracy : 81.00%  
MobileNetV2 Accuracy: 98.50%  
ResNet50 Accuracy : 40.00%

```
[5]: import numpy as np
from sklearn.metrics import confusion_matrix
import seaborn as sns

model = VGG19(weights='imagenet', include_top=False, input_shape=(224,224,3))
model = Sequential([model, GlobalAveragePooling2D(), Dense(128, activation='relu'), Dropout(0.3), Dense(train_data.num_classes, activation='softmax')])
model.compile(optimizer=Adam(1e-4), loss='categorical_crossentropy', metrics=['accuracy'])
model.fit(train_data, validation_data=val_data, epochs=5, verbose=0)
```

```

val_preds = model.predict(val_data)
val_preds = np.argmax(val_preds, axis=1)
val_true = val_data.classes
cm = confusion_matrix(val_true, val_preds)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=val_data.class_indices.keys(),
            yticklabels=val_data.class_indices.keys())
plt.title("Confusion Matrix (VGG19)")
plt.show()

```

7/7

4s 393ms/step

